

Formal Certification of Neural Lyapunov Functions for Transient Stability in Lossy Power Networks

Sehaj Singh

Abstract—A neural Lyapunov function trained against randomly sampled violations cannot be certified, because a certificate quantifies over every point in a region while sampling only ever visits a measure-zero subset, so a network can report a 99.98% empirical satisfaction rate and still fail the Lyapunov conditions on a set that random sampling never lands on. We reproduce the Lyapunov-regularized reinforcement learning of Cui and Zhang for the lossy Kron-reduced New England 39-bus swing model, then show with the neural network verifier CROWN that the as-trained value function violates both the positive-definiteness condition and the Lie-derivative condition through genuine counterexamples that a directed search finds immediately while 200,000 random samples find none. We then close the loop, using the verifier itself as the sampler and feeding its counterexamples back into training, so the same architecture that certifies over nothing becomes certifiable over a real region: on the projected gen-5 slice that matches the authors’ own figure, counterexample-guided training certifies the Lie-derivative condition out to a sublevel radius $\rho = 1.5$ and an exponential decrease rate β up to 3.97, both audited by an independent attack. The certified region collapses as the state dimension grows, which is quantified rather than hidden and which points directly at where the next work has to go.

Index Terms—Neural Lyapunov functions, transient stability, formal verification, CROWN, branch-and-bound, counterexample-guided training.

I. INTRODUCTION

Transient stability of a power network asks whether the generators, after a disturbance, swing back to a synchronous equilibrium rather than losing synchronism. The classical certificate is an energy function, but for a network with resistive losses the conductance coupling enters the swing equations as a term $G_{ij} \cos(\delta_i - \delta_j)$ that is not the gradient of any scalar, so no analytic energy function exists [1]. That single term is why a valid Lyapunov function for lossy frequency dynamics has been an open problem for decades, and it is the entire reason the frequency-control problem is interesting while the voltage-control problem, whose dynamics are linear, is not.

Cui and Zhang sidestep the non-existence result by *learning* a Lyapunov function V_ϕ as a neural network, jointly with a decentralized controller, under a training loss that penalizes violations of the Lyapunov conditions on sampled states [2]. Their method works well empirically, and the last sentence of their paper names the gap it leaves open: an important direction is to verify whether the learned function satisfies the Lyapunov conditions for *all* points in a region, not just on samples. They report the Lie derivative is slightly positive on about 0.1% of sampled points after training. Sampling is

not a proof, because the conditions are universally quantified over a continuous region while any sample set is finite, so the question their paper poses is exactly the question a verifier answers.

This paper answers it. We reproduce their training, we use the neural network verifier CROWN [3] with branch-and-bound over the nonlinear swing dynamics to certify the Lyapunov conditions over a region rather than a sample, and where the as-trained network fails we close the loop with counterexample-guided training. The contribution is not a single verified box, it is the demonstration that the sampling-to-certification gap is real and directed, that it is the same gap in the condition that was supposed to be trivial, and that feeding a sound verifier’s counterexamples back into training is what converts an uncertifiable network into a certifiable one over a real region.

II. RELATED WORK

Certifying a learned Lyapunov function means proving a nonlinear inequality over a continuous region, and the tools differ in what nonlinearities they admit and how they scale. Satisfiability-modulo-theories methods with a δ -complete solver certify neural Lyapunov controllers exactly [4], which is the approach Cui and Zhang cite, but they scale poorly and are typically tractable only for small systems, which is why their own paper notes the limitation. Mixed-integer encodings require piecewise-linear activations and so cannot represent the ELU network used here at all. Sum-of-squares programming requires polynomial dynamics and so cannot represent the sin and cos of the swing equations. Learned-dynamics constructions that build stability in by design [5] avoid verification but change the model class.

Bound propagation with branch-and-bound occupies a different point. CROWN propagates a sound linear relaxation of a general nonlinear computation graph [3], its automatic-differentiation implementation auto_LiRPA bounds arbitrary graphs [6], β -CROWN adds per-neuron branching for completeness [7], and GenBaB extends branch-and-bound to general nonlinearities including sin, cos, and the ELU derivative [8]. The bounds are GPU-parallel across subdomains and, being differentiable, can be trained against [9]. This is the reason CROWN, and not SMT, is the right tool for a lossy swing model with an ELU value function: it handles the exact operators the problem contains, and it hands back a counterexample when it fails rather than only a verdict. Our formulation of the Lie derivative as a boundable graph and the sublevel-set region-of-attraction follow the neural control verification line [10] and its tutorial [11].

III. PROBLEM FORMULATION

A. Lossy swing dynamics

The New England 39-bus system is Kron-reduced to its $n = 10$ generator buses. With state $x = (\delta, \omega) \in \mathbb{R}^{2n}$ the swing dynamics are [2]

$$\dot{\delta}_i = \omega_i, \quad (1)$$

$$M_i \dot{\omega}_i = P_i - D_i \omega_i - u_i(\omega_i) - \sum_j B_{ij} \sin(\delta_i - \delta_j) - \sum_j G_{ij} \cos(\delta_i - \delta_j), \quad (2)$$

where B_{ij} is the susceptance coupling and G_{ij} is the conductance coupling produced by network losses. We match the authors' training dynamics, which integrate $\dot{\delta}_i = \omega_{\text{scale}} \omega_i$ with $\omega_{\text{scale}} = 2\pi$ because their ω is a per-unit frequency in Hz, and which fold the damping $D_i \omega_i$ into the controller, so the certified closed loop is $f(x, u(x))$ with u supplying the damping.

B. Controller and Lyapunov network

The certified controller is the saturated linear droop law $u_i(\omega_i) = \text{sat}(c_i \omega_i)$ that the authors use while training V_ϕ ; it is piecewise linear, so CROWN handles it exactly. The value function $V_\phi(\delta, \omega)$ is a single-hidden-layer network with ELU activations, chosen for a smooth bounded derivative, evaluated on the reference-subtracted coordinates $z = [\delta_i - \delta_1, \omega_i - \omega_1]$ of dimension $2(n-1) = 18$, which removes the rotational symmetry of the swing model so the equilibrium is an isolated point.

C. Lyapunov conditions

With equilibrium x^* and $V^* = V(x^*)$, the conditions are [2]

$$(4a) \quad V(x) > V^* \quad \text{for } x \neq x^*, \quad (3)$$

$$(4b) \quad \dot{V}(x) = \nabla V(x) \cdot f(x, u(x)) < 0 \quad \text{for } x \neq x^*, \quad (4)$$

$$(4c) \quad \dot{V}(x^*) = 0, \quad (5)$$

and Proposition 2 strengthens (4b) to an exponential rate, $\dot{V}(x) \leq -\beta(V(x) - V^*)$. She named (4a) for certification. We certify it because she asked, but (4a) is a property of V alone and never touches the dynamics, so the condition that actually needs the nonlinear swing equations is (4b) and its Proposition-2 form, and that is where the contribution is.

IV. VERIFICATION FORMULATION

A. The Lie derivative as a boundable graph

Condition (4b) requires bounding $\dot{V}(x) = \nabla V(x) \cdot f(x, u(x))$ over a box, which means the gradient of V has to be part of the graph the verifier bounds. For the single-hidden-layer ELU network the gradient is analytic,

$$\nabla_z V = (W_2 \odot \text{ELU}'(W_1 z + b_1)) W_1, \quad \text{ELU}'(a) = e^{-\text{relu}(-a)}, \quad (6)$$

and every operation in it, together with the sin and cos of f , is a natively boundable operator, so we express the whole scalar $F(x) = -\dot{V}(x)$ as one computation graph and bound it with CROWN directly. This specializes the

general-purpose route, in which auto_LiRPA's Jacobian operator expands gradient nodes and GenBaB branches the sin, cos, and ELU nonlinearities [8], [11]; for a one-hidden-layer network the analytic gradient graph is exact and its bounds are tighter, so we use it and keep the Jacobian route available for deeper networks. We also replace the ELU by the identity $\text{ELU}(x) = \text{relu}(x) - \text{relu}(1 - e^{-\text{relu}(-x)})$, which matches `torch.nn.ELU` to 6×10^{-8} and is boundable, since this build of auto_LiRPA has no native ELU bound.

B. Certification and the equilibrium

We certify $F(x) \geq -\epsilon$ over a box B by propagating a sound CROWN lower bound and branching B on its widest dimension where the bound is inconclusive; a nonnegative bound on every leaf verifies the box, and a leaf that cannot be split hands back either a genuine counterexample with $F < -\epsilon$ (the property is false) or an unresolved region (the bound stayed loose). These two outcomes are never conflated. Every verified box is then re-checked by an independent multi-start projected-gradient attack, and a certificate is rejected if the attack finds any $F < -\epsilon$, which guards against an unsound bound.

The Lie derivative is zero at the equilibrium by construction, so $F = -\dot{V}$ is zero there and no region containing x^* certifies (4b) as a strict inequality. We handle this the standard way, certifying on an annulus that excludes a small ball around x^* , and reporting an explicit tolerance ϵ with units where one is used. The region-of-attraction result is a certified sublevel set: we bisect on ρ for the largest $\{x : V(x) \leq \rho\} \cap B$ that verifies, which is forward invariant and is a certified inner approximation of the region of attraction.

C. Soundness by independent cross-check

The whole contribution is soundness, so resting it on one hand-written verifier is a liability, and the independent PGD audit above cannot catch the failure that matters, because an unsound verifier certifies false properties and a PGD attack failing to find a counterexample is exactly what an unsound certificate looks like from the inside. So we cross-check the bounds against a second, independent path that builds the same Lie derivative through auto_LiRPA's JacobianOP, which expands the gradient nodes itself rather than taking our analytic gradient, sharing nothing with the first path on the gradient side. Since both paths bound the same scalar F , a sound lower bound from one must never exceed a sound upper bound from the other, and on 24 boxes spanning the gen-5 region that never happens, with a worst crossing of 0.0 (Table I). The two paths certify the same Proposition-2 rate to the bisection tolerance, and both reject the known-false far region as a violation. The analytic path is the tighter of the two, certifying $\rho = 2.0$ where JacobianOP reaches 1.25 and giving a lower bound of -32.5 where JacobianOP gives -125.0 on the fixed box, so the specialized implementation is a strength that agrees with the general-purpose route rather than a deviation from it.

TABLE I

VERIFIER CROSS-CHECK, SEED 0, GEN-5 SLICE. TWO INDEPENDENT GRADIENT PATHS, SAME CROWN BOUNDS. NO BOX HAS lb FROM ONE PATH ABOVE ub FROM THE OTHER, WHICH IS THE SOUNDNESS CONDITION.

Quantity	analytic gradient	JacobianOP
certified ρ (4b)	2.0	1.25
Prop-2 β	3.94	3.94
far-region verdict	violation	violation
CROWN lower bound, fixed box	-32.5	-125.0
soundness contradictions (24 boxes)	0	

Algorithm 1 Counterexample-guided certification (CEGIS)

```

1: input: trained  $V_\phi$ , target region  $B$ , controller  $u$ 
2: for iteration  $k = 1, 2, \dots$  do
3:    $\rho_k \leftarrow$  largest certifiable sublevel radius on  $B$  (bisection,
     CROWN + branch-and-bound)
4:    $C \leftarrow$  PGD counterexamples maximizing  $\dot{V}$  on the
     frontier just past  $\rho_k$ 
5:   if  $C = \emptyset$  then
6:     return certified  $\rho_k$ 
7:   end if
8:   retrain  $V_\phi$  to drive  $\dot{V} < 0$  on  $C$  and a grid over  $B$ ,
     keeping  $V > V^*$  and  $\dot{V}(x^*) = 0$ 
9: end for

```

V. COUNTEREXAMPLE-GUIDED CERTIFICATION

The as-trained network fails, so we close the loop. Her own Algorithm 1 already collects violating samples by random search, so she has half of this idea; the missing half is that a random sampler visits measure zero while a verifier's counterexamples are sound and concentrate exactly where the property is violated.

VI. CASE STUDY: NEW ENGLAND 39-BUS

All numbers below trace to a per-seed JSON that a script wrote, and every figure regenerates from those JSON files.

A. Reproduction and the sampling gap

The PyTorch port matches the authors' dynamics to 6.4×10^{-7} against an independent transcription of their transfer-matrix algebra. Reproducing their Algorithm 1 gives 99.98% of sampled states with $\dot{V} \leq 0$ and 100% with $V > V^*$, in line with their reported $\approx 99.9\%$. On this same network, uniform random sampling finds a (4b) violation in 0 of 200,000 states, while a directed projected-gradient attack in the same region drives 6.1% of its starts to a violation, with a maximum Lie derivative of +199.9 against -9.6 for the best random sample. The violations are real and the random sampler never lands on them, which is the empirical case for a verifier (Fig. 1).

B. The certification staircase

Branch-and-bound cost grows with input dimension, so we report a staircase rather than claim full-state verification (Table III). On the two-machine gate, a 4-D lossy system, the pipeline certifies (4b) on a real annulus, which is the




Fig. 1. Same network, same region. Random sampling finds no (4b) violation in 200,000 draws, a directed attack finds them at 6.1%. Sampling statistics do not certify.

TABLE II

ROBUSTNESS OVER FIVE SEEDS, MEAN $\pm$ STANDARD DEVIATION. THE RESULT IS NOT A SINGLE LUCKY RUN. QUANTITIES WITH NO SPREAD ARE IDENTICAL ACROSS SEEDS.

Quantity	mean $\pm$ std
E0 sampled decrease-satisfaction	$0.9998 \pm 8 \times 10^{-5}$
E1 random violation rate	$4 \times 10^{-6} \pm 6 \times 10^{-6}$
E1 directed (PGD) violation rate	0.079 ± 0.019
E4 CEGIS certified ρ (4b)	1.5
Prop-2 certified β	3.97
Rung-1 two-bus certified ρ	1.2

correctness check. On the gen-5 projected slice, the 2-D geometry of the authors' own Fig. 2, the as-trained network has a genuine (4b) counterexample with $\dot{V} = +5.22$ on the far region, and CROWN's optimized bound is 2.36 times tighter than interval propagation on the reference box (-54.8 versus -129.5). In the full 20-D state the slice-hardened network genuinely violates (4b) even in a box of radius 0.01, because the counterexample-guided training was scoped to the slice and does not transfer, so the certified full-state radius is reported honestly as the point where a real counterexample appears, not as a verifier timeout.

C. Closing the loop

On the gen-5 slice, counterexample-guided training turns the uncertifiable as-trained network into a certified one. The certifiable sublevel radius goes from 0.0 to 1.5, the full grid maximum, after a single CEGIS iteration, the independent attack then finds a minimum F of +2.43 over the certified region, and no further counterexamples appear through five more iterations (Fig. 3). Proposition 2 certifies an exponential decrease rate β up to 3.97 on the same region, so the result is not a binary verified-or-not but a certified rate. Pushing the region outward, the certified sublevel radius reaches $\rho = 2.0$ before a genuine counterexample appears at $\rho = 2.5$, and across every rung of the staircase every failure the verifier reports is a genuine violation with a counterexample rather than an unresolved loose bound, which says the bounds are tight enough here that not-certified means false, not merely unproven. The same training also fixes (4a): the as-trained network has a genuine positive-definiteness counterexample with $V < V^*$ on the annulus, which 200,000 training samples

TABLE III

CERTIFICATION STAIRCASE, CONDITION (4b), SEED 0. CERTIFIED REGION IS THE SUBLEVEL RADIUS ρ (SLICE) OR BOX RADIUS r (FULL STATE). FAILURES ARE LABELED GENUINE (A COUNTEREXAMPLE EXISTS) OR INCOMPLETENESS (BOUND TOO LOOSE).

Rung	dim	network	result
1: two-bus	4	trained here	certified $\rho = 1.2$, audit $+0.34$
2: gen-5 slice	2	as-trained	genuine viol., $\dot{V} = +5.22$
2: gen-5 slice	2	CEGIS	certified $\rho = 1.5$, audit $+2.43$
3: full state	20	CEGIS	genuine viol. at $r = 0.01$
4: full state	20	CEGIS	genuine viol. at $r = 0.5$, $\dot{V} = +118$

Gen-5 projected slice: sampling trains V to look stable, a verifier finds where it is not




Fig. 2. Gen-5 projected slice. Top, $V - V^*$; bottom, the Lie derivative, red where positive and so violating (4b). Left, the as-trained network with the verifier's counterexample marked; right, the CEGIS network with the certified region drawn in green.

reported satisfied on 100% of draws, and the CEGIS network certifies (4a) there with a dense margin of $+0.097$ (Fig. 2).

D. SMT baseline

The comparison a reviewer expects is against SMT, since Chang, Roohi, and Gao is Cui's reference and her paper dismisses SMT as tractable only for small systems. We encode the identical condition (4b), the same trained network through the same relu/exp ELU identity, for dReal, and check the negation, whether any state in the region has a nonnegative Lie derivative. dReal ships no Windows wheel and its source build needs Bazel and IBEX, so it does not run on the Windows host this work was developed on, the same environment wall the IBM CROWN repository hit, and the encoding is provided to run on Linux where the expected outcome from the SMT scalability literature is that the 4-D two-bus gate closes and the NE39 slice exceeds its budget while CROWN certifies it in seconds. We report the encoding here and do not report a dReal wall-clock we did not measure.

VII. LIMITATIONS

The certificate holds over the stated region and nowhere else. The model is the Kron-reduced lossy swing model that the authors train against, not full AC power flow, and it




Fig. 3. Certified region against CEGIS iteration. Feeding sound counterexamples back into training moves the certified sublevel radius from 0 to 1.5 in one step, and the maximum attack Lie derivative drops below zero and stays there.

omits the explicit damping term, which the controller supplies. The equilibrium is fixed during training and verification, so a certificate that is used to assert stability rather than to regularize would have to account for how the equilibrium moves under load or parameter change, and this one does not. The certified closed loop uses the saturated linear droop law that she trains the Lyapunov function against, so her Algorithm 2 stacked-ReLU RNN controller is not reproduced here and the certificate is for the droop loop rather than for her learned controller, and extending the certificate to the RNN loop is open. Verification cost grows with state dimension, and the counterexample-guided training here was scoped to the projected slice, so the full 20-D result is a genuine violation rather than a certified region, which is the honest boundary of the method as it stands. Everything is simulation, with no hardware.

VIII. CONCLUSION

A learned Lyapunov function that satisfies its conditions on 99.98% of samples can still be falsified by a verifier that quantifies over the whole region, and the same verifier's counterexamples, fed back into training, are what certify a real region on the frequency dynamics that have resisted analytic energy functions for decades. A certified region that grows under counterexample-guided training is the first step toward learned grid controllers that ship with a proof rather than a sampling statistic, and the open problem is that the certified region shrinks as the state dimension grows, which is where the next work, certified training over the full state, has to go.

ACKNOWLEDGMENT

The author thanks Prof. Wenqi Cui for the problem formulation and supervision, and for the New England 39-bus lossy-network data, Kron-reduced with the pg-sync-models toolbox from data in Chow's power system toolbox [12], with droop coefficients obtained by `fmincon`.

REFERENCES

- [1] H.-D. Chiang, "Study of the existence of energy functions for power systems with losses," *IEEE Transactions on Circuits and Systems*, vol. 36, no. 11, pp. 1423–1429, 1989.

- [2] W. Cui and B. Zhang, “Lyapunov-regularized reinforcement learning for power system transient stability,” *IEEE Control Systems Letters*, vol. 6, pp. 974–979, 2022, arXiv:2103.03869.
- [3] H. Zhang, T.-W. Weng, P.-Y. Chen, C.-J. Hsieh, and L. Daniel, “Efficient neural network robustness certification with general activation functions,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2018.
- [4] Y.-C. Chang, N. Roohi, and S. Gao, “Neural lyapunov control,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [5] G. Manek and J. Z. Kolter, “Learning stable deep dynamics models,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [6] K. Xu, Z. Shi, H. Zhang, Y. Wang, K.-W. Chang, M. Huang, B. Kailkhura, X. Lin, and C.-J. Hsieh, “Automatic perturbation analysis for scalable certified robustness and beyond,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [7] S. Wang, H. Zhang, K. Xu, X. Lin, S. Jana, C.-J. Hsieh, and J. Z. Kolter, “Beta-CROWN: Efficient bound propagation with per-neuron split constraints for neural network robustness verification,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2021, arXiv:2103.06624.
- [8] Z. Shi, Q. Jin, J. Z. Kolter, S. Jana, C.-J. Hsieh, and H. Zhang, “Neural network verification with branch-and-bound for general nonlinearities,” *arXiv preprint arXiv:2405.21063*, 2024.
- [9] Z. Shi, H. Li, C.-J. Hsieh, and H. Zhang, “Certified training with branch-and-bound: A case study on lyapunov-stable neural control,” *arXiv preprint arXiv:2411.18235*, 2024.
- [10] L. Yang, H. Dai, Z. Shi, C.-J. Hsieh, R. Tedrake, and H. Zhang, “Lyapunov-stable neural control for state and output feedback: A novel formulation for efficient synthesis and verification,” in *International Conference on Machine Learning (ICML)*, 2024, arXiv:2404.07956.
- [11] H. Zhang *et al.*, “Bridging control with neural network verifier alpha-beta-CROWN: A tutorial,” *arXiv preprint arXiv:2605.26577*, 2026.
- [12] J. H. Chow and K. W. Cheung, “A toolbox for power system dynamics and control engineering education and research,” *IEEE Transactions on Power Systems*, vol. 7, no. 4, pp. 1559–1564, 1992.